

Statistical Computing

Managing data frames with the `dplyr` package

Feng Li

feng.li@cufe.edu.cn

School of Statistics and Mathematics
Central University of Finance and Economics
<http://feng.li>

Data Frames

- *data frame* is a key data structure in statistics and in R.
- one observation per row and each column represents a variable
- we need to have good tools for dealing with them.
- you have seen `subset ()` function and the use of `[` and `$`
- `dplyr` package is designed to mitigate a lot of complex operations for data frames.

The dplyr Package

- by Hadley Wickham of RStudio
- everything dplyr does could already be done with base R, but it *greatly* simplifies existing functionality in R.
- it provides a “grammar” (in particular, verbs) for data manipulation and for operating on data frames.
- the dplyr functions are **very** fast, as many key operations are coded in C++.

dplyr Grammar

Some of the key “verbs” provided by the dplyr package are

- `select`: return a subset of the columns of a data frame, using a flexible notation
- `filter`: extract a subset of rows from a data frame based on logical conditions
- `arrange`: reorder rows of a data frame
- `rename`: rename variables in a data frame
- `mutate`: add new variables/columns or transform existing variables
- `summarise / summarize`: generate summary statistics of different variables in the data frame, possibly within strata
- `%>%`: the “pipe” operator is used to connect multiple verb actions together into a pipeline

Common dplyr Function Properties

All of the functions have a few common characteristics. In particular,

1. The first argument is a data frame.
2. The subsequent arguments describe what to do with the data frame specified in the first argument, and you can refer to columns in the data frame directly without using the \$ operator (just use the column names).
3. The return result of a function is a new data frame.
4. Data frames must be properly formatted and annotated for this to all be useful. In particular, the data must be [tidy](#). In short, there should be one observation per row, and each column should represent a feature or characteristic of that observation.

Installing the dplyr package

```
install.packages("dplyr")
```

After installing the package it is important that you load it into your R session with the `library()` function.

```
library(dplyr)
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

select()

We will use a dataset containing air pollution and temperature data for the [city of Chicago](#) in the U.S.

```
chicago <- readRDS("chicago.rds")
```

```
dim(chicago)
## [1] 6940 8
str(chicago)
## 'data.frame': 6940 obs. of 8 variables:
## $ city : chr "chic" "chic" "chic" "chic" ...
## $ tmpd : num 31.5 33 33 29 32 40 34.5 29 26.5 32.5 ...
## $ dptp : num 31.5 29.9 27.4 28.6 28.9 ...
## $ date : Date, format: "1987-01-01" "1987-01-02" ...
## $ pm25tmean2: num NA NA NA NA NA NA NA NA NA ...
## $ pm10tmean2: num 34 NA 34.2 47 NA ...
## $ o3tmean2 : num 4.25 3.3 3.33 4.38 4.75 ...
## $ no2tmean2 : num 20 23.2 23.8 30.4 30.3 ...
```

Sometimes you may want to use only a couple of variables out of many.

```
names(chicago)[1:3]
## [1] "city" "tmpd" "dptp"
subset <- select(chicago, city:dptp)
head(subset)
## city tmpd dptp
## 1 chic 31.5 31.500
## 2 chic 33.0 29.875
## 3 chic 33.0 27.375
## 4 chic 29.0 28.625
## 5 chic 32.0 28.875
## 6 chic 40.0 35.125
```

Sometimes you may want to drop some variables that are not useful.

```
select(chicago, -(city:dptp))
```

If you wanted to keep every variable that ends with a “2”, we could do

```
subset <- select(chicago, ends_with("2"))
str(subset)
## 'data.frame': 6940 obs. of 4 variables:
## $ pm25tmean2: num NA NA NA NA NA NA NA NA NA ...
## $ pm10tmean2: num 34 NA 34.2 47 NA ...
## $ o3tmean2 : num 4.25 3.3 3.33 4.38 4.75 ...
## $ no2tmean2 : num 20 23.2 23.8 30.4 30.3 ...
```

Or if we wanted to keep every variable that starts with a “d”, we could do

```
subset <- select(chicago, starts_with("d"))
str(subset)

## 'data.frame':    6940 obs. of  2 variables:
##  $ dptp: num  31.5 29.9 27.4 28.6 28.9 ...
##  $ date: Date, format: "1987-01-01" "1987-01-02" ...
```


filter()

The `filter()` function is used to extract subsets of rows from a data frame.

```
chic.f <- filter(chicago, pm25tmean2 > 30)
str(chic.f)
## 'data.frame':    194 obs. of  8 variables:
##  $ city      : chr  "chic" "chic" "chic" "chic" ...
##  $ tmpd      : num  23 28 55 59 57 57 75 61 73 78 ...
##  $ dptp      : num  21.9 25.8 51.3 53.7 52 56 65.8 59 60.3 67.1 ...
##  $ date      : Date, format: "1998-01-17" "1998-01-23" ...
##  $ pm25tmean2: num  38.1 34 39.4 35.4 33.3 ...
##  $ pm10tmean2: num  32.5 38.7 34 28.5 35 ...
##  $ o3tmean2  : num  3.18 1.75 10.79 14.3 20.66 ...
##  $ no2tmean2 : num  25.3 29.4 25.3 31.4 26.8 ...
```

```
summary(chic.f$pm25tmean2)
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  30.05   32.12   35.04   36.63   39.53   61.50
```

We could for example extract the rows where PM2.5 is greater than 30 *and* temperature is greater than 80 degrees Fahrenheit.

```
chic.f <- filter(chicago, pm25tmean2 > 30 & tmpd > 80)
select(chic.f, date, tmpd, pm25tmean2)
##      date tmpd pm25tmean2
## 1 1998-08-23  81  39.60000
## 2 1998-09-06  81  31.50000
## 3 2001-07-20  82  32.30000
## 4 2001-08-01  84  43.70000
## 5 2001-08-08  85  38.83750
## 6 2001-08-09  84  38.20000
## 7 2002-06-20  82  33.00000
## 8 2002-06-23  82  42.50000
## 9 2002-07-08  81  33.10000
##10 2002-07-18  82  38.85000
##11 2003-06-25  82  33.90000
##12 2003-07-04  84  32.90000
##13 2005-06-24  86  31.85714
##14 2005-06-27  82  51.53750
##15 2005-06-28  85  31.20000
##16 2005-07-17  84  32.70000
##17 2005-08-03  84  37.90000
```

arrange()

The `arrange()` function is used to reorder rows of a data frame according to one of the variables/columns.

Here we can order the rows of the data frame by date, so that the first row is the earliest (oldest) observation and the last row is the latest (most recent) observation.

```
chicago <- arrange(chicago, date)
```

We can now check the first few rows

```
head(select(chicago, date, pm25tmean2), 3)
##           date pm25tmean2
## 1 1987-01-01      NA
## 2 1987-01-02      NA
## 3 1987-01-03      NA
```

and the last few rows.

```
tail(select(chicago, date, pm25tmean2), 3)
##           date pm25tmean2
## 6938 2005-12-29    7.45000
## 6939 2005-12-30   15.05714
## 6940 2005-12-31   15.00000
```

Columns can be arranged in descending order too by using the special `desc()` operator.

```
chicago <- arrange(chicago, desc(date))
```

Looking at the first three and last three rows shows the dates in descending order.

```
head(select(chicago, date, pm25tmean2), 3)
##           date pm25tmean2
## 1 2005-12-31   15.00000
## 2 2005-12-30   15.05714
## 3 2005-12-29    7.45000
tail(select(chicago, date, pm25tmean2), 3)
##           date pm25tmean2
## 6938 1987-01-03      NA
## 6939 1987-01-02      NA
## 6940 1987-01-01      NA
```

How would you do this in base R without `dplyr`?

rename()

Renaming a variable in a data frame in R is surprisingly hard to do! The `rename()` function is designed to make this process easier.

Here you can see the names of the first five variables in the `chicago` data frame.

```
head(chicago[, 1:5], 3)
##   city tmpd dptp      date pm25tmean2
## 1 chic   35 30.1 2005-12-31   15.00000
## 2 chic   36 31.0 2005-12-30   15.05714
## 3 chic   35 29.4 2005-12-29    7.45000
```

Now we rename the awkward variable names.

```
chicago <- rename(chicago, dewpoint = dptp, pm25 = pm25tmean2)
head(chicago[, 1:5], 3)
##   city tmpd dewpoint      date   pm25
## 1 chic   35   30.1 2005-12-31 15.00000
## 2 chic   36   31.0 2005-12-30 15.05714
## 3 chic   35   29.4 2005-12-29  7.45000
```

mutate()

The `mutate()` function exists to compute transformations of variables in a data frame.

For example, with air pollution data, we often want to *detrend* the data by subtracting the mean from the data.

```
chicago <- mutate(chicago, pm25detrend = pm25 - mean(pm25, na.rm = TRUE))
head(chicago)
##   city tmpd dewpoint      date    pm25 pm10tmean2  o3tmean2  no2tmean2
## 1 chic   35    30.1 2005-12-31 15.00000      23.5  2.531250  13.25000
## 2 chic   36    31.0 2005-12-30 15.05714      19.2  3.034420  22.80556
## 3 chic   35    29.4 2005-12-29  7.45000      23.5  6.794837  19.97222
## 4 chic   37    34.5 2005-12-28 17.75000      27.5  3.260417  19.28563
## 5 chic   40    33.6 2005-12-27 23.56000      27.0  4.468750  23.50000
## 6 chic   35    29.6 2005-12-26  8.40000       8.5 14.041667  16.81944
##   pm25detrend
## 1   -1.230958
## 2   -1.173815
## 3   -8.780958
## 4    1.519042
## 5    7.329042
## 6   -7.830958
```

group_by()

The `group_by()` function is used to generate summary statistics from the data frame within strata defined by a variable. For example, in this air pollution dataset, you might want to know what the average annual level of PM2.5 is.

First, we can create a year variable using `as.POSIXlt()`.

```
chicago <- mutate(chicago, year = as.POSIXlt(date)$year + 1900)
```

Now we can create a separate data frame that splits the original data frame by year.

```
years <- group_by(chicago, year)
```

Finally, we compute summary statistics for each year in the data frame with the `summarize()` function.

```
summarize(years, pm25 = mean(pm25, na.rm = TRUE), o3 = max(o3tmean2,
  na.rm = TRUE), no2 = median(no2tmean2, na.rm = TRUE))
## # A tibble: 19 x 4
##   year pm25    o3    no2
##   <dbl> <dbl> <dbl> <dbl>
## 1 1987 NaN    63.0  23.5
## 2 1988 NaN    61.7  24.5
## 3 1989 NaN    59.7  26.1
## 4 1990 NaN    52.2  22.6
## 5 1991 NaN    63.1  21.4
## 6 1992 NaN    50.8  24.8
## 7 1993 NaN    44.3  25.8
## 8 1994 NaN    52.2  28.5
## 9 1995 NaN    66.6  27.3
## 10 1996 NaN    58.4  26.4
## 11 1997 NaN    56.5  25.5
## 12 1998 18.3    50.7  24.6
## 13 1999 18.5    57.5  24.7
## 14 2000 16.9    55.8  23.5
## 15 2001 16.9    51.8  25.1
## 16 2002 15.3    54.9  22.7
## 17 2003 15.2    56.2  24.6
## 18 2004 14.6    44.5  23.4
## 19 2005 16.2    58.8  22.6
```

group_by()

In a slightly more complicated example, we might want to know what are the average levels of ozone (o3) and nitrogen dioxide (no2) within quintiles of pm25. A slicker way to do this would be through a regression model, but we can actually do this quickly with `group_by()` and `summarize()`.

First, we can create a categorical variable of pm25 divided into quintiles.

```
qq <- quantile(chicago$pm25, seq(0, 1, 0.2), na.rm = TRUE)
chicago <- mutate(chicago, pm25.quint = cut(pm25, qq))
```

Now we can group the data frame by the `pm25.quint` variable.

```
quint <- group_by(chicago, pm25.quint)
```

Finally, we can compute the mean of o3 and no2 within quintiles of pm25.

```
summarize(quint, o3 = mean(o3tmean2, na.rm = TRUE), no2 = mean(no2tmean2,
  na.rm = TRUE))
## # A tibble: 6 x 3
##   pm25.quint    o3    no2
##   <fct>        <dbl> <dbl>
## 1 (1.7,8.7]    21.7  18.0
## 2 (8.7,12.4]  20.4  22.1
## 3 (12.4,16.7] 20.7  24.4
## 4 (16.7,22.6] 19.9  27.3
## 5 (22.6,61.5] 20.3  29.6
## 6 <NA>        18.8  25.8
```




The pipeline operator `%>%` is very handy for stringing together multiple `dplyr` functions in a sequence of operations.

```
third(second(first(x)))
```

This nesting is not a natural way to think about a sequence of operations. The `%>%` operator allows you to string operations in a left-to-right fashion, i.e.

```
first(x) %>% second %>% third
```



Take the example that we just did in the last section where we computed the mean of o3 and no2 within quintiles of pm25. There we had to

1. create a new variable `pm25.quint`
2. split the data frame by that new variable
3. compute the mean of o3 and no2 in the sub-groups defined by `pm25.quint`

That can be done with the following sequence in a single R expression.

```
mutate(chicago, pm25.quint = cut(pm25, qq)) %>% group_by(pm25.quint) %>%
  summarize(o3 = mean(o3tmean2, na.rm = TRUE), no2 = mean(no2tmean2,
    na.rm = TRUE))
## # A tibble: 6 x 3
##   pm25.quint    o3    no2
##   <fct>      <dbl> <dbl>
## 1 (1.7,8.7]    21.7  18.0
## 2 (8.7,12.4]  20.4  22.1
## 3 (12.4,16.7] 20.7  24.4
## 4 (16.7,22.6] 19.9  27.3
## 5 (22.6,61.5] 20.3  29.6
## 6 <NA>       18.8  25.8
```


Summary

The `dplyr` package provides a concise set of operations for managing data frames. With these functions we can do a number of complex operations in just a few lines of code. In particular, we can often conduct the beginnings of an exploratory analysis with the powerful combination of `group_by()` and `summarize()`.

- `dplyr` can work with other data frame “backends” such as SQL databases. There is an SQL interface for relational databases via the `DBI` package
- `dplyr` can be integrated with the `data.table` package for large fast tables

The `dplyr` package is handy way to both simplify and speed up your data frame management code. It’s rare that you get such a combination at the same time!

Lab Session

dplyr

You'll be working with the `airquality` in the R package `datasets`. Bear in mind `%>%`.

1. Please return all the rows where `Temp` is larger than 80 and `Month` is after May.
2. Please add a new column that displays the temperature in Celsius.
3. Calculate the mean temperature in each month.
4. Remove all the data corresponding to `Month = 5`, group the data by month, and then find the mean of the temperature each month.

References

Chapter 13 of the book “R programming for data science”.